



كلية الملك عبد الله الثاني  
لتكنولوجيا المعلومات  
KING ABDULLAH II SCHOOL OF  
INFORMATION TECHNOLOGY

# **Machine Learning and Neural Networks**

## **Time Series Prediction and Classification**

Team members:

Jana Godieh 0223457

Islam Alodat 0226399

Kasandra Haddad 0227097

Shahed Alemian 2221209

**Under supervision of: Dr. Ali \_ Alrweidan**

# 1.0 Introduction:

## 1.1 Datasets

### 1.1.1 Breast Cancer Wisconsin Dataset:

This dataset is available on Kaggle and was originally created by Dr. William H. Wolberg from the University of Wisconsin-Madison. The dataset is used for binary classification tasks, specifically to predict whether a breast mass is **malignant (cancerous)** or **benign (non-cancerous)** based on features derived from digitized images of fine needle aspirates (FNA) of breast masses.

Overview of the Dataset:

1. **Number of Attributes:** 32
  - 30 numeric features (predictors)
  - 1 ID column
  - 1 target column (diagnosis)
2. **Target Variable:** diagnosis
  - M = Malignant
  - B = Benign

Features:

The dataset contains the following features, which are computed from digitized images of breast mass FNA samples. These features describe characteristics of the cell nuclei present in the image:

1. **ID:** Unique identifier for each sample.
2. **Diagnosis:** Target variable (M = Malignant, B = Benign).
3. **Ten Real-Valued Features for Each Nucleus:**
  - **Radius:** Mean distance from the center to points on the perimeter.
  - **Texture:** Standard deviation of gray-scale values.
  - **Perimeter:** Perimeter of the nucleus.
  - **Area:** Area of the nucleus.
  - **Smoothness:** Local variation in radius lengths.
  - **Compactness:**  $\text{Perimeter}^2 / \text{area} - 1.0$ .
  - **Concavity:** Severity of concave portions of the contour.
  - **Concave points:** Number of concave portions of the contour.
  - **Symmetry:** Symmetry of the nucleus.

- **Fractal dimension:** "Coastline approximation" - 1.

For each of these features, the dataset provides the following statistics:

- mean: Mean value.
- se: Standard error.
- worst: Worst (largest) value

We will use this data to build a predictive model aimed at classifying the tumors as either benign or malignant based on these features.

### 1 .1.2 Hourly Energy Consumption dataset:

PJM Interconnection LLC (PJM) is a regional transmission organization (RTO) in the United States. It is part of the Eastern Interconnection grid operating an electric transmission system.

The hourly power consumption dataset contains over 10 years of hourly energy consumption data from PJM's website and is in megawatts (MW). We obtained the Dataset from Kaggle.

Since the data includes energy consumption on an hourly basis, it allows for the analysis of patterns and trends over extended periods. For the purpose of this task, we will choose ( AEP\_hourly.csv) and applied time series forecasting .

The dataset includes the following features:

- **Datetime:** The timestamp for each observation (hourly intervals).
- **Energy Consumption:** The amount of energy consumed (in MW)

## 1.2 Overview of tasks:

### 1.2.1 Breast Cancer Classification:

The task for this dataset was to implement the following models:

1. **Multi-Layer Perceptron (MLP):** Design and train a fully connected feedforward neural network.
2. **Support Vector Machine (SVM):** Evaluate performance using linear and kernel-based approaches.
3. **Ensemble Learning Technique:** Implement Random Forest.

### 1.2.2 Time Series for Hourly Energy Consumption:

The task for this dataset was to implement the following models:

1. **Echo State Network (ESN):** Focus on its efficiency in sequential data processing.
2. **Long Short-Term Memory (LSTM):** Train and evaluate using a traditional LSTM architecture.
3. **Bi-Directional LSTM (Bi-LSTM):** Apply a bi-directional structure to capture past and future context.

## 1.3 Objectives

### 1.3.1 objectives for breast cancer classification:

The goal of tumor classification is to determine whether it is benign or malignant, which helps in making appropriate treatment decisions. By improving diagnostic accuracy and reducing reliance on human expertise, errors can be avoided, and the diagnosis can be more accurate using the automated model. It also helps save time and costs, especially in crowded healthcare centers, by providing a faster and more efficient tool.

The model assists doctors in making better treatment decisions based on predictions derived from the data. It also aids in analyzing medical data and understanding the relationship between tumor characteristics and its status,

enhancing the screening and diagnostic process. Additionally, it helps in early prediction, improving treatment and recovery chances.

By using machine learning techniques, large volumes of medical data can be handled, extracting patterns that are difficult for humans to detect. It also enables monitoring disease progression and the impact of various treatments, providing a helpful tool for doctors in making accurate clinical decisions. Overall, the goal is to improve diagnostic and treatment processes by building an accurate and effective model for tumor classification.

### 1.3.2 The objectives of hourly energy consumption:

Analyzing consumption patterns and understanding how they change over time is of paramount importance to inform policy decisions or infrastructure planning. It helps improve energy efficiency by identifying high consumption periods and peak times. It also contributes to energy management planning and better resource allocation. Additionally, it helps respond to fluctuations in energy demand, enhancing the ability to meet energy needs during critical times. For example, a utility company might use this dataset to forecast peak demand periods and adjust generation accordingly. Researchers might analyze the impact of weather conditions (e.g., temperature) on energy consumption to help in those efforts.

## 2.Data preprocessing steps taken:

### 2.1 The data preprocessing in classification:

We followed the following steps to preprocess the data.

#### 1. Read the Data

```
[2]: df = pd.read_csv("cancer.csv")
df
```

```
[2]:
```

	id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	compactness_mean	concavity_mean	concave points_mean	...	texture
0	842302	M	17.99	10.38	122.80	1001.0	0.11840	0.27760	0.30010	0.14710	...	
1	842517	M	20.57	17.77	132.90	1326.0	0.08474	0.07864	0.08690	0.07017	...	
2	84300903	M	19.69	21.25	130.00	1203.0	0.10960	0.15990	0.19740	0.12790	...	
3	84348301	M	11.42	20.38	77.58	386.1	0.14250	0.28390	0.24140	0.10520	...	
4	84358402	M	20.29	14.34	135.10	1297.0	0.10030	0.13280	0.19800	0.10430	...	
...	...	...	...	...	...	...	...	...	...	...	...	
564	926424	M	21.56	22.39	142.00	1479.0	0.11100	0.11590	0.24390	0.13890	...	
565	926682	M	20.13	28.25	131.20	1261.0	0.09780	0.10340	0.14400	0.09791	...	

**2.Detection of missing values**, where all columns are checked to determine if they contain missing values

```
[3]: missing_values = df.isnull().sum()
missing_values
```

```
[3]: id                                0
diagnosis                             0
radius_mean                           0
texture_mean                           0
perimeter_mean                         0
area_mean                             0
smoothness_mean                       0
compactness_mean                      0
concavity_mean                        0
concave points_mean                   0
symmetry_mean                         0
fractal_dimension_mean                0
radius_se                             0
texture_se                             0
perimeter_se                          0
area_se                               0
smoothness_se                         0
compactness_se                        0
concavity_se                          0
concave points_se                     0
symmetry_se                           0
fractal_dimension_se                  0
radius_worst                          0
texture_worst                         0
perimeter_worst                       0
area_worst                            0
smoothness_worst                      0
compactness_worst                     0
concavity_worst                       0
concave points_worst                  0
symmetry_worst                        0
fractal_dimension_worst                0
Unnamed: 32                           569
dtype: int64
```

3. **Cleaning:** The purpose of this step is to clean the dataset by eliminating columns that do not add any analytical value or do not impact the desired outcome. We removed the ID column and column 32.  
the ID column does not provide any meaningful information or predictive power for the classification task (predicting whether a tumor is malignant or benign). Including it might cause the model to mistakenly interpret the ID as a feature.

```
[4]: df = df.drop(['id', 'Unnamed: 32'], axis=1)
```

4. Converting the diagnosis into numerical values so that it can be processed by the model.

```
[5]: le = LabelEncoder()  
df['diagnosis'] = le.fit_transform(df['diagnosis'])
```

5. Separating the target feature (the diagnosis) into a separate variable (y)

```
[6]: y = df['diagnosis']  
x = df.drop('diagnosis', axis = 1)  
print(x.shape)  
print(y.shape)
```

```
(569, 30)
```

```
(569,)
```


6. Quantify the class distribution in the dataset. We did this to count the number of instances for each class in the target variable (diagnosis)

```
[9]: print(y.value_counts())
```

```
diagnosis  
0      357  
1      212  
Name: count, dtype: int64
```

7. Visualizes the class distribution of the target variable

```
[10]: y.value_counts().plot.pie(autopct = '%.2f')  
plt.show()
```



8. Splitting the data into 30% test and 70% train

```
[11]: x_train,x_test,y_train,y_test = train_test_split(x,y,test_size = 0.3,random_state = 42)  
print(f'train samples:{len(x_train)}')  
print(f'test samples:{len(x_test)}')
```

train samples:398  
test samples:171

9. **Feature selection:** We used a heat map to guide our **feature selection** by highlighting redundant features that provide similar information. This step is very important to remove redundant features. This can improve model efficiency.



10. Identify highly correlated features based on a predefined threshold ( $> 0.85$ ).

```
[13]: col_corr = set()
      for i in range(len(corr.columns)):
          for j in range(i):
              if corr.iloc[i,j] > 0.85:
                  colname = corr.columns[i]
                  col_corr.add(colname)
      col_corr
```

```
[13]: {'area_mean',
      'area_se',
      'area_worst',
      'compactness_worst',
      'concave points_mean',
      'concave points_worst',
      'concavity_mean',
      'concavity_worst',
      'perimeter_mean',
      'perimeter_se',
      'perimeter_worst',
      'radius_worst',
      'texture_worst'}
```

11. Remove the highly correlated features

```
[14]: x_train = x_train.drop(col_corr, axis = 1)
      x_test = x_test.drop(col_corr,axis = 1)
```

12. We used SMOTE to fix the imbalance in the data. We did not use an under-sampling technique because we do not want to lose parts of the data. To avoid data leakage we only used SMOTE on the training data set

```
[15]: print(y_train.value_counts())
```

```
diagnosis
0      249
1      149
Name: count, dtype: int64
```


```
[16]: sm = SMOTE(random_state =42)
      x_res,y_res = sm.fit_resample(x_train,y_train)
```

And this was the result

```
[17]: print(y_res.value_counts())
```

```
diagnosis
0      249
1      249
Name: count, dtype: int64
```

```
[18]: y_res.value_counts().plot.pie(autopct = '%.2f')
plt.show()
```



## 2.2 Data preprocessing in time series:

### 2.2.1 Echo state preprocessing

## 1. Read the data

```
[2]: df = pd.read_csv('AEP_hourly.csv')
```

```
[3]: df.head()
```

```
[3]:
```

	Datetime	AEP_MW
0	2004-12-31 01:00:00	13478.0
1	2004-12-31 02:00:00	12865.0
2	2004-12-31 03:00:00	12577.0
3	2004-12-31 04:00:00	12517.0
4	2004-12-31 05:00:00	12670.0

## 2. Check for missing values

```
[10]: print(df.isnull().sum())
```

```
Datetime    0
AEP_MW      0
dtype: int64
```

## 3. Sort the data

```
[11]: df = df.sort_values(by='Datetime')
```

```
[12]: df.head()
```

```
[12]:
```

	Datetime	AEP_MW
<b>2183</b>	2004-10-01 01:00:00	12379.0
<b>2184</b>	2004-10-01 02:00:00	11935.0
<b>2185</b>	2004-10-01 03:00:00	11692.0
<b>2186</b>	2004-10-01 04:00:00	11597.0
<b>2187</b>	2004-10-01 05:00:00	11681.0

## 4. Split the data into 70% train and 30% test

```
[17]: train_size = int(len(df) * 0.7)
      train, test = df[:train_size], df[train_size:]
```

## 5. Creating input-output pairs with a one-step lag

```
[19]: x_train = train['AEP_MW'].values[:-1].reshape(-1, 1)
      y_train = train['AEP_MW'].values[1:].reshape(-1, 1)
      x_test = test['AEP_MW'].values[:-1].reshape(-1, 1)
      y_test = test['AEP_MW'].values[1:].reshape(-1, 1)
```

## 6. Scaling the data so all the data has the same magnitude

```
[20]: scaler = MinMaxScaler()
      x_train_scaled = scaler.fit_transform(x_train)
      y_train_scaled = scaler.fit_transform(y_train)
      x_test_scaled = scaler.transform(x_test)
      y_test_scaled = scaler.transform(y_test)
```

## 2.2.2. LSTM and BI-LSTM preprocessing

### 1. Read the data

```
[33]: data = df['AEP_MW'].values.reshape(-1, 1)
```

### 2. Scale the data

```
[34]: data_scaled = scaler.fit_transform(data)
```

### 3. Split the data into 80% train 10% test 10% validation.

```
[35]: train_size = int(len(data_scaled) * 0.8)
      validation_size = int(len(data_scaled) * 0.10)
      test_size = len(data_scaled) - train_size - validation_size
```

```
[36]: train_set = data_scaled[: train_size]
      validation_set = data_scaled[train_size : train_size + validation_size]
      test_set = data_scaled[train_size + validation_size :]
```

### 4. Preparing time-series data. We create input-output pairs using a look-back window, enabling the model to learn temporal dependencies and patterns in

the data.

```
[37]: def create_sequences(data, look_back):
      x, y = [], []
      for i in range(len(data) - look_back):
          x.append(data[i:i + look_back])
          y.append(data[i + look_back])
      return np.array(x), np.array(y)

[38]: look_back = 24
      x_train, y_train = create_sequences(train_set , look_back)
      x_val, y_val = create_sequences(validation_set , look_back)
      x_test, y_test = create_sequences(test_set , look_back)
```

## 3. Methodology

### 3.1 Multilayer perceptron

#### Model Architecture:

The model was implemented using `tensorflow.keras.models Sequential` and `tensorflow.keras.layers`'s `Dense` and `Dropout`

#### 1. Input Layer:

- The input shape matches the train data's shape (`x_res_scaled.shape[1]`)

#### 2. Hidden Layers:

- The number of hidden layers is determined by the `n_layers` parameter.
- Each hidden layer uses the **ReLU (Rectified Linear Unit)** activation function (`activation='relu'`).
- The number of units in each hidden layer decreases by a factor of 2 for each subsequent layer:
- **Dropout** is applied after each hidden layer to prevent overfitting. The dropout rate is controlled by the `dropout_rate` parameter.

#### 3. Output Layer:

- The output layer has **1 unit** with a **sigmoid activation function**, so that the output is a probability between 0 and 1

#### 4. Compilation:

- The model is compiled using the **Adam optimizer**.
- The loss function is **binary\_crossentropy**
- The evaluation metric is **accuracy**.

## Hyperparameters:

The model's hyperparameters are tuned using **GridSearchCV**. The values explored were:

1. **n\_layers:**
  - Number of hidden layers in the neural network.
  - Values: [1, 2, 3]
2. **units:**
  - Number of neurons in the first hidden layer. The number of neurons in subsequent layers decreases by a factor of 2.
  - Values: [32, 64, 128]
3. **dropout\_rate:**
  - Dropout rate applied after each hidden layer to prevent overfitting.
  - Values: [0.3, 0.4, 0.5]
4. **batch\_size:**
  - Number of samples processed before the model is updated.
  - Values: [16, 32]
5. **epochs:**
  - Number of times the model trains on the entire dataset.
  - Values: [40, 50, 60]

## 3.2 Support Vector Machine (SVM)

### Model Architecture:

We used `sklearn.svm`'s SVC model.

### Hyperparameters:

The SVM model is tuned using **RandomizedSearchCV** and the hyperparameter values used were:

1. **C:**
  - **Regularization Parameter:** Controls the trade-off between maximizing the margin and minimizing classification errors.
  - **Values:** Sampled from a uniform distribution between 0.1 and 10.

2. **gamma:**

- **Kernel Coefficient:** Defines the influence of a single training example (only for RBF and polynomial kernels).
- **Values:** Sampled from a uniform distribution between 0.01 and 1.

3. **kernel:**

- **Kernel Function:** Determines the type of transformation applied to the data.
- **Values:** ['linear', 'rbf', 'poly'].

4. **degree:**

- **Degree of Polynomial Kernel:** Specifies the degree of the polynomial function (only for the polynomial kernel).
- **Values:** [2, 3, 4].

### 3.3 Random forest classifier

#### Model Architecture:

The random forest was implemented using **sklearn.ensemble's RandomForestClassifier**. The key components of the model are:

#### Key Components:

We used the library provided by sklearn (RandomForestClassifier)

#### Hyperparameters:

The Random Forest model is tuned using **GridSearchCV**, and multiple different scorers were tested using a for loop. The Hyperparameters explored were:

1. **n\_estimators:**

- **Controls the number of trees in the forest.**
- **Values:** np.arange(50, 150, 10).

2. **max\_depth:**

- **Maximum depth of each tree.**
- **Values:** [2, 3, 4, 5, 6].

3. **min\_samples\_split:**

- Controls the minimum number of samples required to split an internal node.
  - Values: [20, 30, 40, 50].
4. **min\_samples\_leaf:**
- Controls the minimum number of samples required to be at a leaf node.
  - Values: [10, 15, 20, 25].

The scorers explored were:

1. **Accuracy:** `make_scorer(accuracy_score)`
2. **Precision:** `make_scorer(precision_score, average='macro')`
3. **Recall:** `make_scorer(recall_score, average='macro')`
4. **F1:** `make_scorer(f1_score, average='macro')`
5. **ROC\_AUC:** `make_scorer(roc_auc_score)`

## 3.4 LSTM

### Model Architecture:

The LSTM model is implemented using **tensorflow.keras.models Sequential** and **tensorflow.keras.layers LSTM , Dense and Dropout** .It consists of the following layers:

1. **Input Layer:**
  - The input shape is `(x_train.shape[1], x_train.shape[2])`, where:
    - `x_train.shape[1]` is the **sequence length** (number of time steps in each input sequence we decided to create sequences of 24 hours or one day).
    - `x_train.shape[2]` is the **number of features** in each time step.
2. **First LSTM Layer:**
  - **Neurons:** 50
  - **Activation:** `tanh` , (`activation = "tanh"`)
  - **Return Sequences:** `True`, (`return_sequences = True`)  
meaning the layer returns the output for each time step in the sequence



### 3. First Dropout Layer:

- **Dropout Rate:** 0.2
- Dropout is applied to prevent overfitting.

### 4. Second LSTM Layer:

- **Units:** 50
- **Activation:** tanh
- **Return Sequences:** False, meaning the layer returns only the output of the last time step (final prediction).

### 5. Dropout Layer:

- **Dropout Rate:** 0.2
- Dropout is applied to prevent overfitting.

### 6. Output Layer:

- **Dense Layer:** 1 unit (uses a linear activation).
- This layer outputs a single value, which is the predicted value for the next time step.

## Model Compilation:

- **Optimizer:** adam
- **Loss Function:** mean\_squared\_error (MSE)

## Hyperparameters:

1. **Number of LSTM Units:** 50 in both LSTM layers.
2. **Activation Function:** tanh for LSTM layers.
3. **Dropout Rate:** 0.2.
4. **Batch Size:** 32.
5. **Epochs:** 20.

## 3.5 Echo State Network

### Model Architecture:

The model was implemented using `reservoirpy.nodes` Reservoir and Ridge

The ESN model consists of two main components:

#### 1. Reservoir:

- A large, fixed, and randomly connected recurrent network of neurons.

- The reservoir captures the temporal dynamics of the input data.
- 2. **Readout Layer:**
  - A linear output layer (Ridge Regression) that maps the reservoir states to the target output.

## Key Components:

1. **Reservoir:**
  - **Size:** 100 neurons
  - **Spectral Radius (sr):** 0.9, which controls the scaling of the reservoir's weight matrix to ensure stability.
  - **Connectivity (rc\_connectivity):** 0.08, meaning only 8% of the reservoir's internal connections are active (sparse connectivity).
  - **Input Scaling (input\_scaling):** 0.9, which scales the input weights to the reservoir.
  - **Seed:** 42, for reproducibility of the random initialization.
2. **Readout Layer:**
  - **Ridge Regression:**
    - **Ridge Parameter (ridge):** 1e-6, which controls the strength of regularization.
3. **Connecting the reservoir and the ridge:**  
(esn\_model = reservoir >> ridge)  
The >> operator connects the reservoir to the readout layer, creating the ESN model.

## Hyperparameters:

1. **Reservoir:**
  - **Size:** 100 neurons.
  - **Spectral Radius (sr):** 0.9.
  - **Connectivity (rc\_connectivity):** 0.08.
  - **Input Scaling (input\_scaling):** 0.9.
  - **Seed:** 42.
2. **Readout Layer:**
  - **Ridge Parameter (ridge):** 1e-6.

## Workflow:

### 1. Model Initialization:

- The reservoir and readout layer are initialized with the specified hyperparameters.
- The `>>` operator connects the reservoir to the readout layer, creating the ESN model.

### 2. Training:

- The model is trained on the scaled training data (`x_train_scaled`, `y_train_scaled`).
- During training, the reservoir states are computed for the input data, and the readout layer is trained to map these states to the target output.

### 3. Prediction:

- The trained model is used to predict the target values for the scaled test data (`x_test_scaled`).

### 4. Inverse Scaling:

- The predicted and actual values are inverse-transformed using a scaler to convert them back to the original scale.

### 5. Evaluation:

- The model's performance is evaluated using metrics like RMSE, MAE, and  $R^2$  on both the scaled and original data.

## 3.6 BI-LSTM

### Model Architecture:

The BI-LSTM model is implemented using `tensorflow.keras.models Sequential` and `tensorflow.keras.layers Bidirectional`, `LSTM`, `Dense` and `Dropout`. It consists of the following layers:

#### 1. Input Layer:

- The input shape is (`x_train.shape[1]`, `x_train.shape[2]`), where:
  - `x_train.shape[1]` is the **sequence length** (number of time steps in each input sequence we chose 24 which represents one day of data).
  - `x_train.shape[2]` is the **number of features** in each time step.

## 2. First BI-LSTM Layer:

- **Bidirectional:** Processes the sequence in both directions.
- **Neurons:** 50
- **Activation:** tanh (activation = 'tanh')
- **Return Sequences:** True, meaning the layer returns the output for each time step in the sequence

## 3. Dropout Layer:

- **Dropout Rate:** 0.2
- Dropout is applied to prevent overfitting.

## 4. Second Bi-LSTM Layer:

- **Bidirectional:** Processes the sequence in both directions.
- **Units:** 50
- **Activation:** tanh
- **Return Sequences:** True, meaning the layer returns the output for each time step in the sequence

## 5. Dropout Layer:

- **Dropout Rate:** 0.2
- Another dropout layer is applied after the second BI-LSTM layer.

## 6. Output Layer:

- **Dense Layer:** 1 unit (uses a linear activation).
- This layer outputs a single value, which is the predicted value for the next time step.

## Model Compilation:

- **Optimizer:** adam
- **Loss Function:** mean\_squared\_error (MSE)

## Hyperparameters:

1. **Number of BI-LSTM Neurons:** 50 in both BI-LSTM layers.
2. **Activation Function:** tanh
3. **Dropout Rate:** 0.2 (applied after each BI-LSTM layer).
4. **Batch Size:** 32 (number of samples processed before the model is updated).

5. **Epochs:** 20 (number of times the model trains on the entire dataset).

## 4. Experiment Setup:

Training environment:

For all 6 modules implemented we used the anaconda distribution of jupyter notebook as our programming environment

**In the classification task**

We used the following libraries for **preprocessing**: numpy,pandas ,sklearn.model\_selection's train\_test\_split ,imblearn.over\_sampling's SMOTE ,sklearn.preprocessing's LabelEncoder and StandardScaler

We used the following libraries for **visualization**:

Matplotlib.pyplot and seaborn

Specific libraries for each model:

**Multilayer perceptron:**

tensorflow

tensorflow.keras.models Sequential

tensorflow.keras.layers Dense,Dropout

scikeras.wrappers KerasClassifier

sklearn.model\_selection's GridSearchCV

sklearn.metrics

### **RandomForest :**

Sklearn.ensemble RandomForestClassifier,sklearn.model\_selection's  
GridSearchCV ,sklearn.metrics

### **Support vector machine:**

Sklearn's svm and sklearn.svm's SVC and scipy.stats's uniform and  
sklearn.model\_selection's RandomizedSearchCV

### **In the Timeseries task**

We used the following libraries for **preprocessing**: numpy,pandas , and  
sklearn.preprocessing's MinMaxScaler

We used the following libraries for **visualization**:

Matplotlib.pyplot ,seaborn,plotly.graph\_objects

Specific libraries for each model:

### **Echo state network:**

Reservoirpy.nodes's Reservoir and ridge and sklearn.metrics  
mean\_squared\_error,mean\_absolute\_error and r2\_score

### **LSTM and BI-LSTM:**

tensorflow

tensorflow.keras.models Sequential

tensorflow.keras.layers Dense,Dropout,LSTM,Input,Bidirectional

sklearn.metrics mean\_squared\_error,mean\_absolute\_error and r2\_score

## Implementation Details:

### 4.1 Multilayer perceptron

Our first attempt at implementing our multi-layer perceptron we used 3 layers with 64,32,1 neuron in each and we didn't use any Drop out layers,

when we fit the model and tested its performance, we found out our model was overfitting so to overcome this we decided to use 2 dropout layers with a 0.5 drop out rate which solved our overfitting problem but we still wanted to improve our performance and that's when we used sklearn's GridSearchCV

to be able to use it we needed to import a wrapper function, so we used kerasClassifier provided by scikeras.wrappers

and then we created a function (called create\_nn) that will create MLP's when fed the parameters (n\_layers , units, dropout\_rate) and we chose relu to be our activation as per standard practice and we fixed it as a parameter

the function does the following:

- 1.create a sequential model
- 2.add a dense layer that takes as input a matrix that matches the shape of our data (input\_shape=(x\_res\_scaled.shape[1],))
3. add a drop out layer
4. loop through the passed number of layers and calculate the number of units for each new layer (by dividing the last number by 2)
  - 4.1 It will create the layer with the calculated number of units
  - 4.2 add a drop out layer

5. after the loop is done it will add one final layer that will generate the final output

6. compile the model using adam as an optimizer and using binary\_crossentropy as the loss function

After implementing the function, we used gridsearchcv with the previously specified parameter grid and we also used 5 folds of cross validation

So, the final workflow was:

**1. Model Creation:**

- The create\_nn function defines the neural network architecture based on the hyperparameters passed to it.

**2. Hyperparameter Tuning:**

- GridSearchCV trains and evaluates the model for each combination of hyperparameters using 5-fold cross-validation.

**3. Best Model Selection:**

- After the search, the best hyperparameters and the corresponding model are selected.

**4. Evaluation:**

- The best model is used to make predictions on the test set (x\_test\_scaled).

We then ran grid search on our train data and using it found the best parameters which were

**Best Parameters: {'batch\_size': 32, 'epochs': 60, 'model\_\_dropout\_rate': 0.4, 'model\_\_n\_layers': 3, 'model\_\_units': 128}**

Which gave us our best model



## 4.2 Random Forest

From the first time we implemented the model we decided to use gridSearchCV and a loop of different scorers to try to find the best possible model

Our implementation loops through the scorers and runs gridSearchCV using the RandomForestClassifier provided by sklearn, the parameter grid specified, the current scorer of the loop and uses 10-fold validation each scorer's best result is stored in an array to see the best model using the scorer

Our initial tries used this parameter grid and 5-fold validation:

```
parameters = {  
    'n_estimators': [10,20,30,40,50],  
    'max_depth': [2,3,4,5,6,7],  
    'min_samples_split': [10,20,30, 50],  
    'min_samples_leaf': [20,30,40,50]  
}
```

This gave us 94 accuracy and 86 recall which didn't satisfy us especially on recall as a missed diagnosis would cause a lot of harm, so we tried another set of parameters

```
parameters = {  
    'n_estimators': [10,50,100,500],  
    'max_depth': [2, 4, 8, 16, None],  
    'min_samples_split': [2,10,20,30, 50],  
    'min_samples_leaf': [1, 2, 4]  
}
```

Which resulted in these accuracy and recall results (94, 92) but we wanted better performance especially on recall as this is used for medical diagnosis, so we decided to lift our cross-validation parameter to 10 folds and change the ranges of our parameter grid to become this

```
parameters = {
    'n_estimators': np.arange(50, 150, 10),
    'max_depth': [2, 3, 4, 5, None],
    'min_samples_split': [20, 30, 40, 50],
    'min_samples_leaf': [10, 15, 20, 25]
}
```

This grid helped in achieving higher performance with 96 accuracy and 98 recall but this grid caused our model to overfit, so we decided to limit the max\_depth parameter to help improve overfitting.

```
parameters = {
    'n_estimators': np.arange(50, 150, 10),
    'max_depth': [2, 3, 4, 5, 6],
    'min_samples_split': [20, 30, 40, 50],
    'min_samples_leaf': [10, 15, 20, 25]
}
```

This helped us achieve the best results with no overfitting.

### Training:

The model was trained on our training data (x\_res\_scaled, y\_res), And our best parameters were:

Best Hyperparameters: {'max\_depth': 6, 'min\_samples\_leaf': 10, 'min\_samples\_split': 20, 'n\_estimators': 130}

### Prediction:

Our best model was used to predict the target values for the test data (x\_test\_scaled).

## 4.3 Support Vector Machine

The first time we implemented the model, we used the following parameters:

**C = 1.0:** This value controls the flexibility of the model.

**gamma = 'scale':** Adjusts the spatial influence of data points.

**kernel = 'rbf':** A radial basis function kernel was chosen to create nonlinear decision boundaries. When we trained the model with these parameters, and tested the model, it showed low accuracy. We discovered that it was possible to improve the model's accuracy by increasing the hyperparameters.

To solve this problem, we used RandomizedSearchCV to increase the number of hyperparameter values. The parameter distribution was set as follows:

**C:** Values between 0.1 and 10 were selected using a uniform distribution.

**gamma:** Values between 0.01 and 1 were explored using a uniform distribution.

**kernel:** Different kernel types (linear, rbf, poly) were tested.

**degree:** Values 2, 3, and 4 were tested for the polynomial kernel

We performed the random search using 50 random samples with 5-layer cross-validation. The evaluation metric used was accuracy.

## **Results:**

After completing the search, the best hyperparameters for the model were:

**Best Parameters: {'C': 1.5092422497476266, 'degree': 4, 'gamma': 0.1752669390630025, 'kernel': 'linear'}**

The optimized model showed significant improvement compared to the first model, with accuracy increasing from a low baseline to a much higher value.

We used the optimized model to make predictions on the test set, and we evaluated the performance using:

**Confusion Matrix:** which showed a noticeable reduction in errors compared to the initial model.

**Classification Report**, providing metrics such as accuracy, recall, and F1-score for each class.

**Overall Accuracy**, which significantly increased after hyperparameter tuning.

## 4.4 Echo State

Our first attempts at implementing the echo state network we used reserviourpy's ESN library, but the results weren't ideal, so we attempted to run gridsearchCV using this library to achieve better results, but ESN isn't compatible with sklearn's gridsearchCV as there is no wrapper class available. our next experiment was changing the library, so we tried pyESN's ESN library and it also didn't support gridsearchCV and the results also weren't ideal.

So, we turned back to reserviourpy and tried their reserviour and ridge models.

Which gave us the best results out of all 3 choices we tried, so we started tuning the hyperparameters manually based on the results on the test dataset among the many values we tried were these:

```
reservoir = Reservoir(  
    1000,  
    sr=0.85,  
    rc_connectivity=0.2,  
    input_scaling=0.7,  
    seed=42  
)  
ridge = Ridge(ridge=1e-3)  
reservoir = Reservoir(  
    2000,  
    sr=0.8,  
    rc_connectivity=0.15,  
    input_scaling=0.6,  
    seed=42  
)  
ridge = Ridge(ridge=1e-3)
```

```
reservoir = Reservoir(  
    2000,  
    sr=0.75,  
    rc_connectivity=0.1,  
    input_scaling=0.5,  
    seed=42  
)  
ridge = Ridge(ridge=5e-3)
```

But our best results were from these parameters:

```
reservoir = Reservoir(  
    100,  
    sr=0.9,  
    rc_connectivity=0.08,  
    input_scaling=0.9,  
    seed=42  
)  
ridge = Ridge(ridge=1e-6)  
esn_model = reservoir >> ridge
```

After finding the best parameters we trained the model on our training data

## 4.5 LSTM

To find the best number of layers and nodes we tried observing what other programmers usually used when training the LSTM model and we found out that almost everyone uses 2 layers and from 30 to 60 nodes in each layer and from 10 to 30 epochs of training, we also noticed people used relu and tanh activation functions, thus we decided to manually try the combinations to see which gave us the best results

To help us in picking the best results we used the validation set to tune the hyperparameters.


The model is trained using the following parameters:

- **Training Data:** x\_train (input sequences) and y\_train (target values).

- **Validation Data:**  $x_{val}$  (validation input sequences) and  $y_{val}$  (validation target values).

We tried 40 for the first layer and 40 in the second, We tried 40 and 30 ,we also t tried 30 and 20 ,50 and 40 and then 50 and 50 which gave us the best results after that we tried different number of epochs we tried 10 and 20 and 25 and 22 and 30 and 40

40 caused very bad overfitting.



So, we tried 20 which gave us the best results.

All the layers and trials we did used tanh as it is the standard in LSTM.

## 4.6 BI-LSTM

In implementing the BI-LSTM we tried using the same parameters we used in our LSTM model as they gave us impressive results and after training the model on the parameters, we found from our LSTM model we also got very good results.

The model is trained using the following parameters:

Training Data: x\_train (input sequences) and y\_train (target values).

Validation Data: x\_val (validation input sequences) and y\_val (validation target values).

So, we decided to use them too, but we tried multiple training epoches like what we did in our LSTM but with a tighter range we tried 20 ,22,25,30,40 and judged them based on our validation results.

After multiple tries, we got the best results from 20 epoches.

## 5. Results and Analysis:

The classification models results were:

### Multilayer Perceptron



The classification report on test data:

	precision	recall	f1-score	support
0	0.98	0.96	0.97	108
1	0.94	0.97	0.95	63
accuracy			0.96	171
macro avg	0.96	0.97	0.96	171
weighted avg	0.97	0.96	0.97	171

Accuracy: 0.9649122807017544

Comparing the results with the baseline accuracy provided by UCI for this dataset which was (87.413 – 96.503) our model performance was very impressive.

Visualizing the results:





## Random Forest

The classification report on test data:

```
Test Set Performance:
              precision    recall  f1-score   support

     0           0.99       0.94       0.96       108
     1           0.90       0.98       0.94        63

 accuracy          0.95       171
 macro avg         0.94       0.96       0.95       171
 weighted avg      0.96       0.95       0.95       171
```

The classification report on train data:

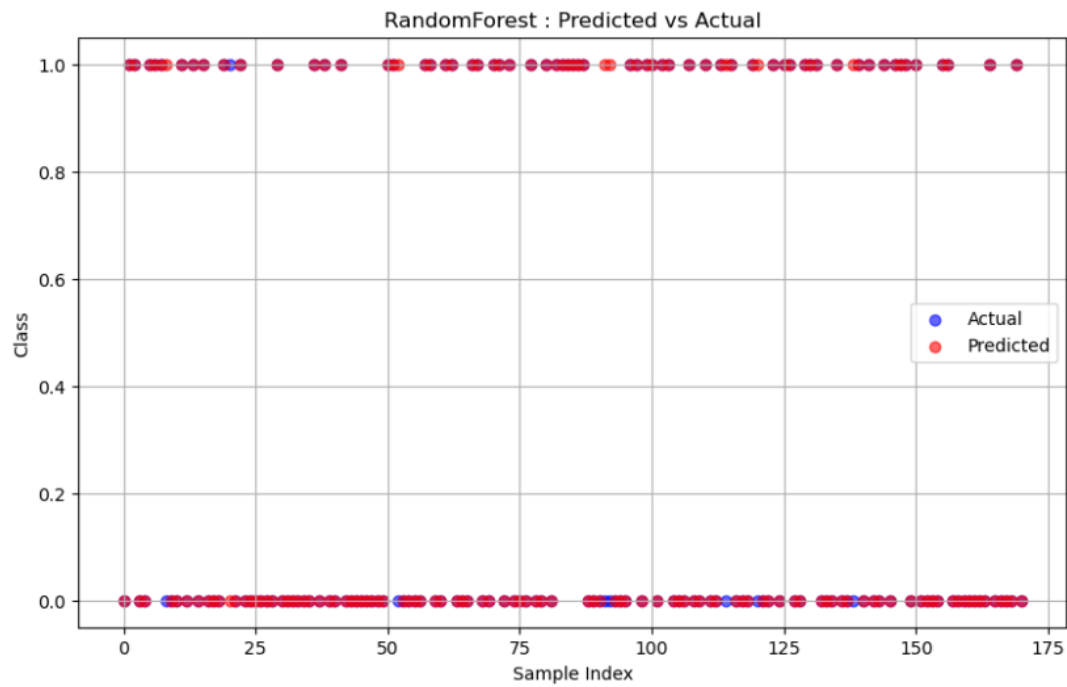

```
              precision    recall  f1-score   support

     0           0.97       0.98       0.97       249
     1           0.98       0.97       0.97       249

 accuracy          0.97       498
 macro avg         0.97       0.97       0.97       498
 weighted avg      0.97       0.97       0.97       498
```

Comparing the results with the baseline accuracy provided by UCI for this dataset which was (95– 100) our model performance was on the lower end but our goal here was to ensure that recall is as high as possible for the malignant detection which we achieved with a 0.98 score.

Visualizing the results:



## Support Vector Machine (SVM)



The classification report on test data:

	precision	recall	f1-score	support
0	0.97	0.97	0.97	108
1	0.95	0.95	0.95	63
accuracy			0.96	171
macro avg	0.96	0.96	0.96	171
weighted avg	0.96	0.96	0.96	171


Accuracy: 0.9649122807017544

Comparing the results with the baseline accuracy provided by UCI for this dataset which was (90.210– 97.902) our model performance was on the higher range.

Visualizing the results:



Comparing all 3 models together graphically:



Based on the previous results our best performing model was:

Accuracy-wise our multi-layer perceptron and svm models gave the same results.

Recall-wise our random forest model performed the best.

Overall our MLP model had the best results over all metrics

The timeseries models' results were:

## Echo state Network

Scaled test results using RMSE, MSE,  $R^2$ :

```
Scaled Test RMSE: 0.014443004862413763
```

```
Scaled Test MAE: 0.010728367608010589
```

```
Scaled Test R^2: 0.9913666874005821
```

Test results after inverse transformation:

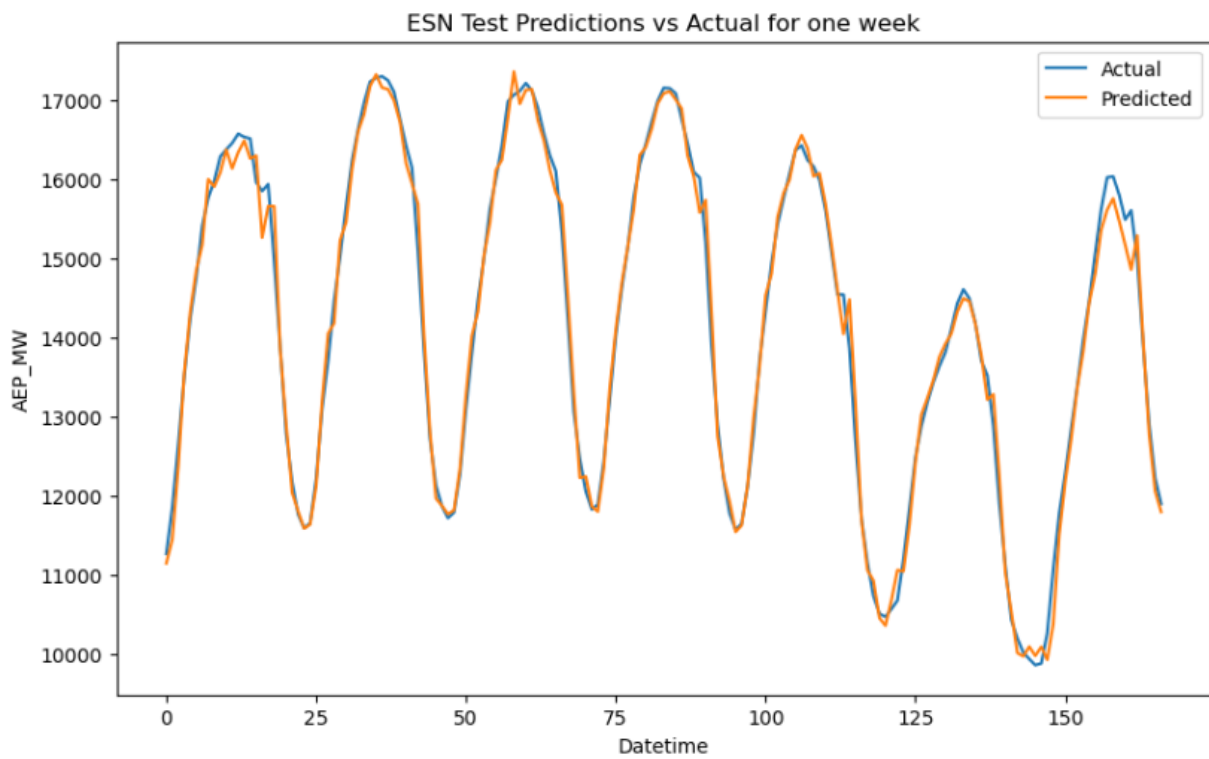
---



Test RMSE: 231.46359592504297

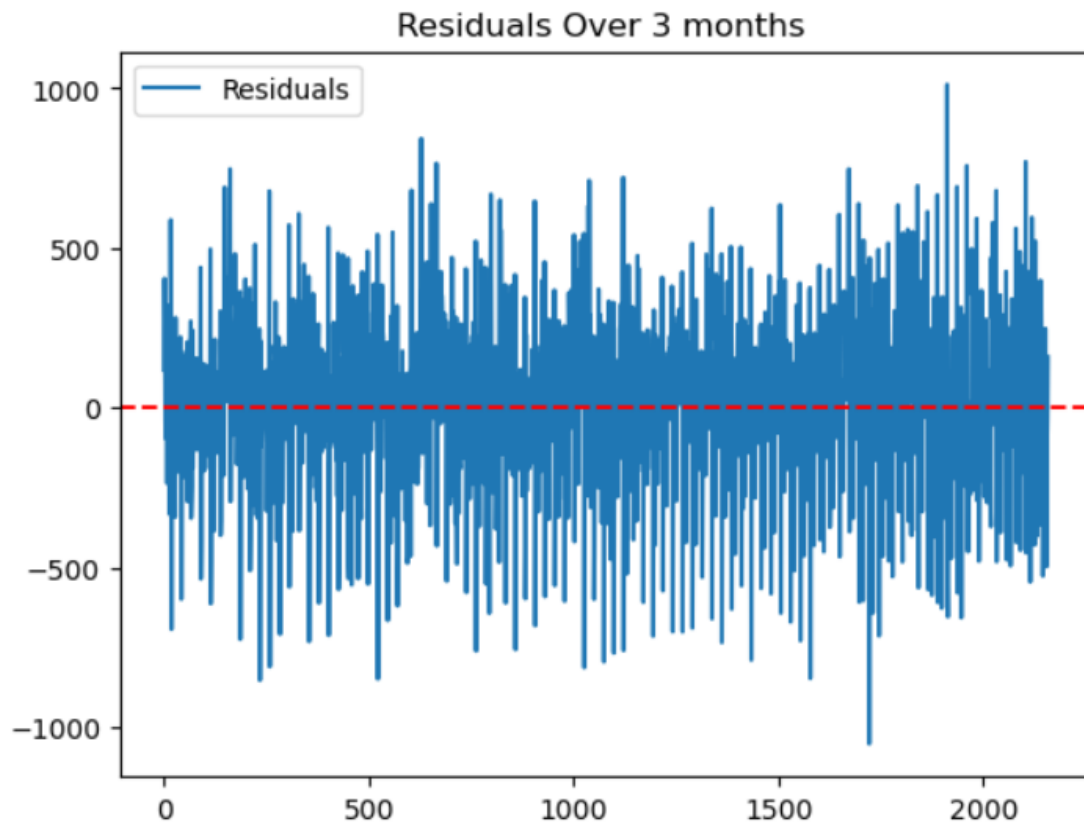
Test MAE: 171.9328192859777

Test  $R^2$ : 0.9913666874005821

Visualizing the results:








## LSTM

Validation loss compared with train loss:



Test results with RMSE, MAE, R<sup>2</sup> after inverse transformation:





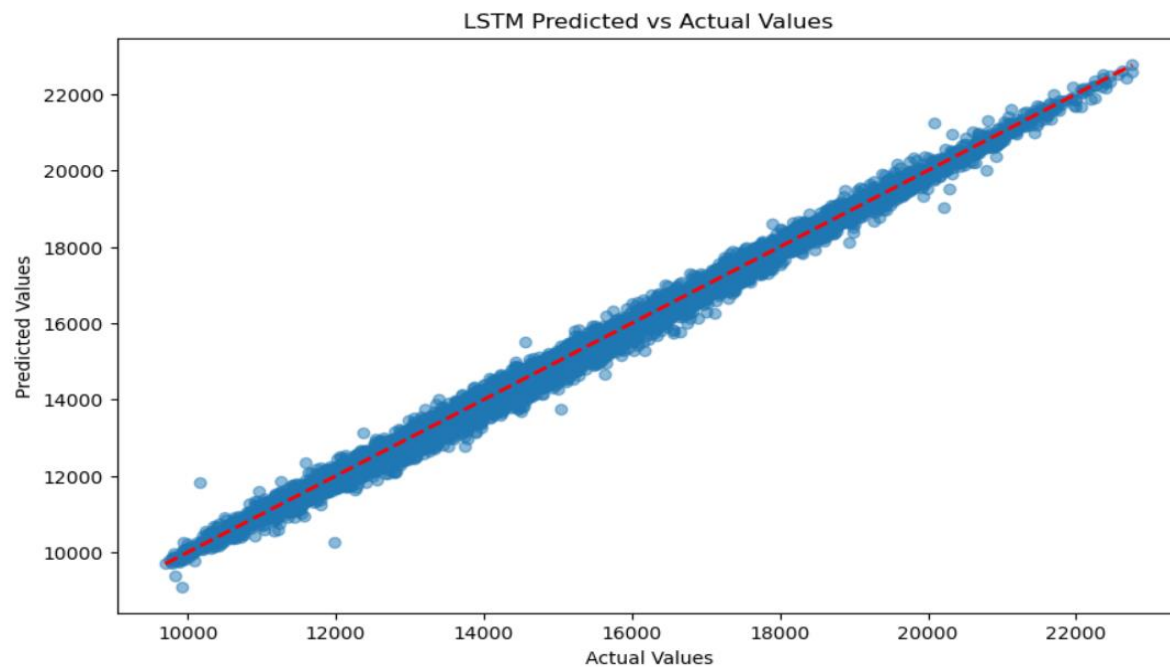
Test LSTM RMSE: 178.87416002070714

Test LSTM MAE: 131.93691119025735

Test LSTM R<sup>2</sup>: 0.9946996277242551


Visualizing the results:





## BI-LSTM

Validation loss compared with train loss:

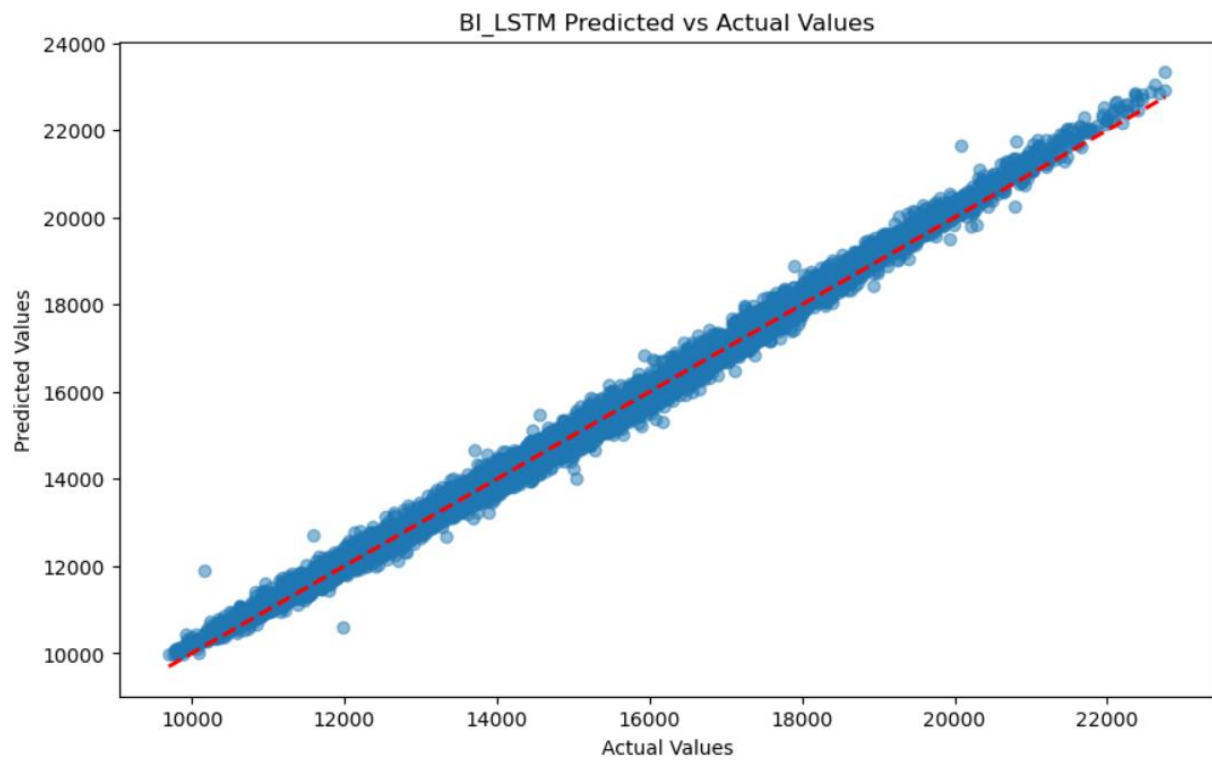
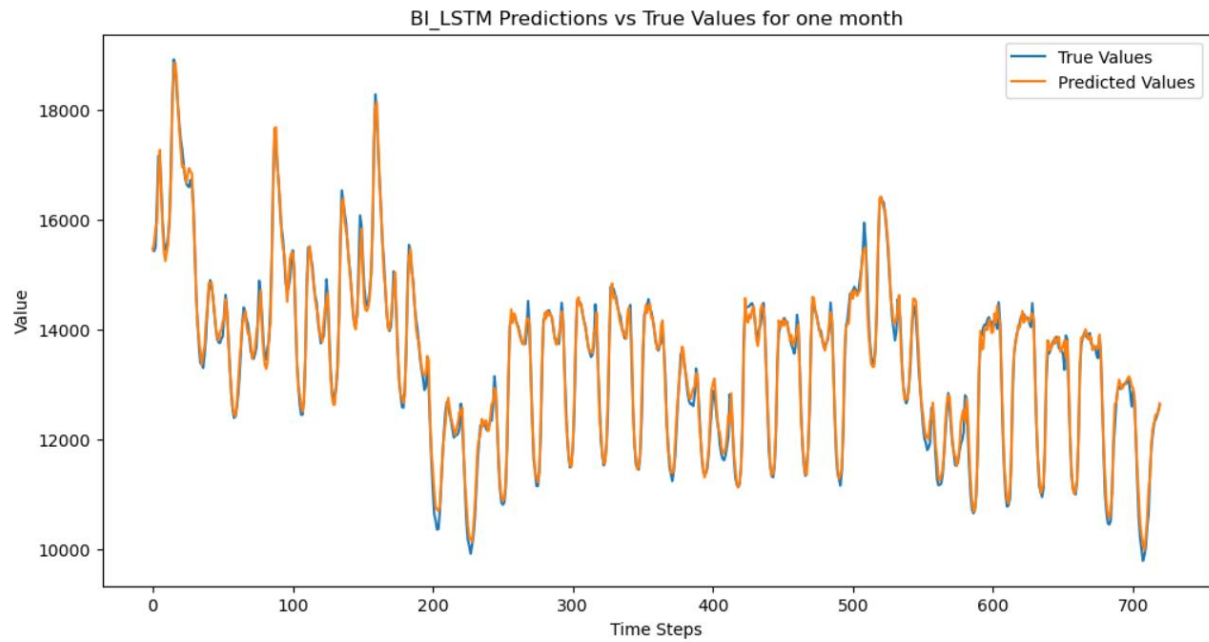


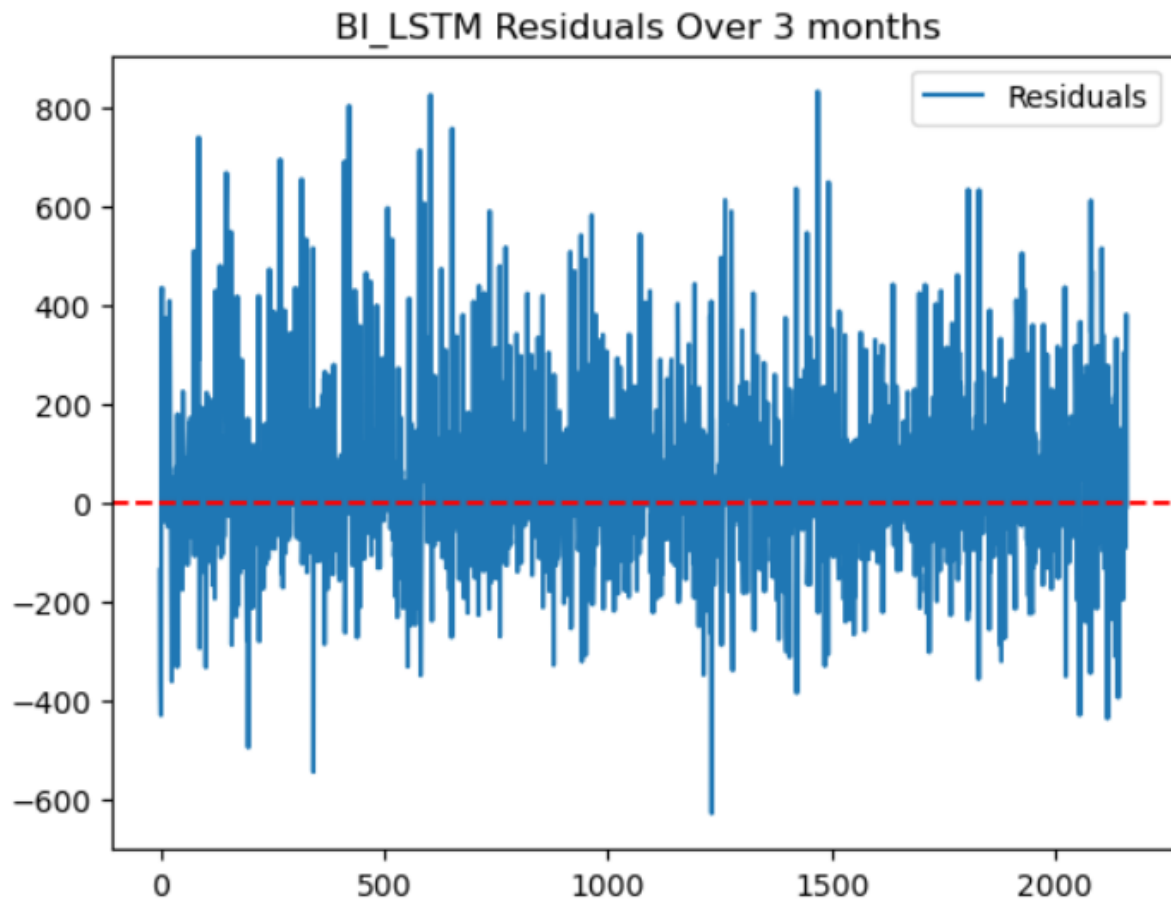
Test results with RMSE, MAE, R^2 after inverse transformation:

---

```
Test BI_LSTM RMSE: 184.78396648819484
Test BI_LSTM MAE: 139.75436607406223
Test BI_LSTM R^2: 0.9943436050033698
```

Visualizing the results:





Based on the previous results our best performing model was the LSTM

Commentary on residuals plots:

Our actual values are in the 12,000-19,000 range, so our residuals ranges through all three plots represent relatively small errors (about 2-3% of the values) compared to the actual data ranges this is also further supported by our  $R^2$  scores for all 3 models

## 6. Conclusion and Recommendations

In this project, we explored two distinct tasks: **Breast Cancer Classification** and **Hourly Energy Consumption Time Series Prediction**. For each task, we implemented and evaluated multiple machine learning models, using different techniques and methods focusing on their performance, accuracy, and ability to generalize unseen data.

### **Breast Cancer Classification:**

- **Multi-Layer Perceptron (MLP):** We implemented our best model using tensorflow and tuned the hyperparameters with gridsearchCV achieving impressive results with an accuracy of 96% on the test set, outperforming the baseline accuracy provided by UCI (87.413% - 96.503%). The model demonstrated strong performance in classifying both benign and malignant tumors, with high precision, recall, and F1-scores.
- **Random Forest:** We implemented our best model using sklearn RandomForestClassifier and tuned the hyperparameters with gridsearchCV and different scorers achieving a test accuracy of 95%, with a particularly high recall of 98% for malignant cases, which was a key objective given the medical implications of false negatives.
- **Support Vector Machine (SVM):** We implemented our best model using sklearn.svm's SVC and tuned the hyperparameters with randomsearchCV achieving a test accuracy of 96%, placing it on the higher end of the baseline accuracy range (90.210% - 97.902%).

Overall, the **MLP** model was the best-performing model for the breast cancer classification task, with the great performance across all metrics.

### **Hourly Energy Consumption Time Series Prediction:**

- **Echo State Network (ESN):** We implemented our best model using reservoirpy's reservoir and ridge and tuned the hyperparameters manually based on the test results and graphical representation of results achieving an RMSE of 231.4 and a MAE of 171.9 and  $R^2$  of 0.99.
- **Long Short-Term Memory (LSTM):** We implemented our best model using Tensorflow and tuned the hyperparameters manually based on the validation results and graphical representation of results. The LSTM model outperformed the ESN with an RMSE of 178.8 and MAE of 131.9 and  $R^2$  of 0.99. The LSTM model was able to capture temporal dependencies effectively, resulting in accurate predictions of energy consumption trends.
- **Bidirectional LSTM (Bi-LSTM):** We implemented our best model using Tensorflow and tuned the hyperparameters manually based on the validation results and graphical representation of results. The Bi-LSTM achieved an RMSE of 184.7 and MAE of 139.7 and  $R^2$  of 0.99.
- **Recommendations for Future Work**
  1. **Feature Engineering:** For the time series task, including supplementary features like weather data, economic indicators or holidays and festive seasons could improve the model's ability to capture external factors affecting energy consumption.
  2. **Advanced Architectures:** Exploring more advanced neural network architectures could improve the performance of classification and time series tasks.
  3. **Detecting anomalies in energy consumption:** any further work in this topic can focus on detecting anomalies in energy consumption data, this could provide valuable insights into identifying unusual patterns or potential issues in the energy grid.

## Final Thoughts

This project successfully demonstrated the effectiveness of various machine learning models in both classification and time series prediction tasks. The results highlight the potential of machine learning in various fields as demonstrated by its ability to improve diagnostic accuracy in healthcare and optimize energy consumption forecasting.

## 7. References:

1. **Breast Cancer Wisconsin (Diagnostic) Data Set**

<https://www.kaggle.com/datasets/uciml/breast-cancer-wisconsin-data>

2. **Hourly Energy Consumption**

<https://www.kaggle.com/datasets/robikscube/hourly-energy-consumption>